

Charge what it's worth.


Value-based dynamic pricing for agentic products.

I've seen this problem from the inside.

GenieLabs AI	Head of Product
Trigo	AI Product Manager
SGC Knowledge Transfer	Director of Product

Eight years in product, the last four in AI. Pricing is the decision product people rarely make. This sits on it.

The whole market is changing how it charges.



Gartner says agentic AI “breaks the link between user growth and revenue growth” and puts **\$234B of app spend at risk by 2030**.

ICONIQ, State of AI, three waves, n approx 300 software companies each. Gartner, July 2026.

This market was born out of SaaS.


Intercom, then Fin	\$0.99 per outcome
Zendesk, then AI agents	per automated resolution
Replit, then Agent	per checkpoint, by effort

The same companies, and the same pricing DNA they arrived with: a flat rate per unit .

intercom.com/pricing, zendesk.com/pricing, docs.replit.com/billing. Read 2026-09-01.

What changed is the value.

SAAS



Value per user barely moves, and cost moves with it.


— Value — Cost

AGENTS



Value swings per request, and cost swings on its own.

Agents get more autonomous. The band widens.



Solid is METR's measured time horizon, doubling every 4.2 months.

METR, Measuring AI Ability to Complete Long Software Tasks, Horizon v1.1. Dashed is that trend extended.

Same product, two requests, one price list.

"Make the button gray"	0.50 credits
"Add authentication"	1.20 credits
"Build a landing page with images"	2.00 credits

Lovable's own published price list, 0.50 to 2.00. One is a colour. One is whether the product can ship. **The price moves 4x. The value does not.**

docs.lovable.dev/user-guides/messaging-limits. Read 2026-09-01.

*A professional prices
from two numbers.
An agent prices from **one**.*

A different price for every request.

01

The value

What this buyer accepts for this request. What is asked, and who is asking.

$\text{price} = \max(\text{floor}, \text{argmax}_p [p \times P(\text{accept} | p)])$

A different price per request, set before the run, explained by the rule that set it .

02

The floor

Cost p90 plus the target margin. This is the part that is built.

03

The decision

Where to land between them, learned from every accept and every reject.

A platform can build the scorer, not the prior.

A new customer has no acceptance history, so the first price has to come from somewhere. It comes from customers with the same declared attributes: deal size, close rate, vertical. That makes cross-customer data a condition, not an asset .



The floor is built. Nothing else is.

Cost prediction demo	running, local
Value scorer	not built
The decision engine	not built
Customers, pilots, revenue	zero

Next proof: a backtest of a differential price list over one customer's own transaction history.

Three introductions, not money yet.

- WHO Companies selling a unit of output whose quality varies.
- ASK Three introductions this week, and their transaction history.
- THEN A date now for the second meeting.

The category already exists. Thoma Bravo paid \$1.4B for PROS in December 2025, and Zilliant sells this same decision to industrial B2B. Nobody sells it to agentic products.

Thoma Bravo, completion of the PROS acquisition, December 2025. docs.zilliant.com.